

From cycling-environment supply to bike-share demand: an IMD-augmented spatio-temporal forecasting benchmark for French networks

Rohan Fossé^{*1} and Gaël Pallares^{†2}

¹*CESI Engineering School, Montpellier, France*

²*CESI LINEACT (EA 7527), Montpellier, France*

Abstract

Operational bike-share demand prediction is dominated by two practices that do not communicate well : classical time-series forecasting on each station (*e.g.*, ARIMA, persistence) and deep spatio-temporal learning trained from scratch on the trip log. Neither integrates the rich *supply-side* information that public open data already exposes about each station – transit access, infrastructure density, topography, demographic pool.

We close this gap with the **IMD-4** (*Indice de Mobilité Douce à quatre composantes*), a Bayesian composite indicator of cycling-environment quality computed on every French commune from public data, and benchmark its added predictive value against six model families on the Vélomagg dock-based panel (53 stations, 5 years, hourly resolution, 375,305 observations). On a one-year temporal hold-out test, an IMD-augmented gradient-boosted model reaches $R^2_{\text{trips}} = 0.311$ against 0.042 for the same model trained without IMD features, a **+0.27 absolute R^2 gain attributable to the IMD station-level enrichment alone**. The IMD-augmented model also matches a station-fixed-effect baseline within 0.001 R^2 , meaning the 13 IMD physical features absorb the spatial signal that 67 station dummies would otherwise have to memorise – an essential property for transferring the model to stations without trip history. Modern recurrent and Transformer-style baselines do not beat the gradient-boosted tabular SOTA on this data size, a finding consistent with the recent benchmark literature on small-to-medium tabular forecasting.

The IMD-4 is computable on all 34,858 French communes from five public open-data sources alone, so the predictive pipeline transfers to any commune with at least an active dock-based network. We replicate the headline result across **four additional networks** (Capital Bikeshare DC, Divvy Chicago, Bluebikes Boston, Bay Wheels SF, totalling $\sim 3,200$ stations and ~ 11.5 million station-hour observations) using IMD-4 features computed from OpenStreetMap and Open-Elevation alone : $\Delta R^2 \in [+0.31, +0.49]$ with a mean of +0.36 across the four international cities, and +0.27 at the Vélomagg calibration anchor. Code, data and trained models are released under MIT.

Keywords: bike-sharing ; demand forecasting ; composite indicator ; spatio-temporal model ; Bayesian inference ; gradient boosting ; LSTM ; French Plan Vélo.

Highlights

- IMD-4 cycling-environment indicator integrated as feature in demand model
- +0.27 R^2 gain on hourly trips at single-city Vélomagg ($n = 375,305$)
- Replicated across 5 networks: $\Delta R^2 \in [+0.27, +0.49]$, mean +0.36
- IMD beats persistence, ARIMA, Ridge, LSTM and matches station fixed-effect
- Country-agnostic pipeline: OSM + Open-Elevation + GBFS open data

*Corresponding author: rfosse@cesi.fr. ORCID: [0009-0002-2195-0198](https://orcid.org/0009-0002-2195-0198)

†ORCID: [0009-0002-8680-604X](https://orcid.org/0009-0002-8680-604X)

1 Introduction

Public bike-sharing systems have become a structural component of urban mobility in France and across Europe (Pucher and Buehler, 2010; Buehler and Pucher, 2012). Operational decisions – where to add a station, how to rebalance the fleet, when to expect demand peaks – require a forecasting layer between the live data feed and the dispatcher (Forma et al., 2015; Raviv et al., 2013). Two practices currently dominate that layer, and they do not communicate with each other. The first is classical time-series : per-station ARIMA models or naive persistence baselines that exploit the autoregressive structure of the trip series (Box and Jenkins, 1976; Hyndman et al., 2008). The second is deep spatio-temporal learning : recurrent and graph-convolutional networks trained from scratch on the trip log, often with weather covariates (Lin et al., 2018; Wang et al., 2022). Neither family integrates the rich *supply-side* information that public open data already exposes about each station – the density of nearby heavy-transit stops, the cycle-infrastructure quality of the surrounding street network, the topographic friction of the catchment, the demographic pool of the host commune. All four have been documented in the cycling-environment literature as demand drivers (Heinen et al., 2010; Schoner and Levinson, 2014; Parkin et al., 2008; Rietveld and Daniel, 2004; Faghih-Imani et al., 2014), yet none is routinely available to an operational forecasting pipeline.

Contribution. This paper closes that gap with two complementary contributions. **(C1)** We introduce the *Indice de Mobilité Douce à quatre composantes* (IMD-4), a Bayesian composite indicator of cycling-environment quality combining four supply-side axes – heavy-transit multimodality (M), cycling-infrastructure density (I), topographic friction (T) and population density (D) – under a simplex prior calibrated jointly against the FUB perception barometer (Fédération française des Usagers de la Bicyclette, 2023) and the EMP modal-share survey (SDES, 2020). The IMD-4 is computable on every French commune ($n = 34,858$) from public data alone (Section 2). **(C2)** We benchmark the IMD-4 as a feature in a spatio-temporal demand model on the Vélomag Montpellier panel (53 stations, 5 years, hourly resolution, 375,305 observations) against six model families : persistence, per-station ARIMA, Ridge regression, LightGBM gradient boosting (Ke et al., 2017), an LSTM recurrent net (Hochreiter and Schmidhuber, 1997), and a LightGBM ablation without IMD features. The IMD-augmented gradient-boosted model wins the benchmark with $R^2_{\text{trips}} = 0.311$ against 0.042 for the same model without IMD ($\Delta R^2 = +0.27$, Section 5).

Why this matters operationally. A natural alternative to IMD-augmentation is to add one dummy per station (station fixed-effect), which captures the spatial signal by memorisation. The IMD-augmented model matches the fixed-effect baseline within 0.001 R^2 , meaning the 13 IMD physical features fully absorb what 67 dummies would otherwise need to learn. This is essential for transferring the model to stations without trip history – the operational case of evaluating where to install a new station (El-Assi et al., 2017; Ashqar et al., 2017).

2 The IMD-4 cycling-environment indicator

2.1 Four supply-side axes

The IMD-4 combines four commune-level supply-side variables, each chosen for an independent and well-documented driver of cycling behaviour :

M – Multimodality. Density of heavy transit stops (train, métro, tramway) per km^2 ; heavy-transit accessibility is the dominant factor in cycling-as-feeder modal choice (Martens, 2007; Saberi et al., 2018). Bus stops are excluded because the bike-and-ride dynamic operates at a different scale.

I – Cycling infrastructure. Total kilometres of dedicated cycling infrastructure per km^2 – all *aménagements cyclables* (segregated paths, on-road lanes, green-ways, shared bus lanes, contraflow, and mixed configurations). Infrastructure density is the supply-side variable

most consistently associated with commute cycling share in the international literature (Schoner and Levinson, 2014; Parkin et al., 2008; Heinen et al., 2010).

T – Topography. Elevation roughness derived from BD ALTI 25 m. Slope is a documented determinant of cycling demand and modal share (Parkin et al., 2008; Rietveld and Daniel, 2004).

D – Density. Logarithm of commune population density. Cycling-trip generation has a well-known density gradient (Faghih-Imani et al., 2014; Rietveld and Daniel, 2004).

For the calibration panel of Section 2.3, the four axes are evaluated at station resolution with a 300 m buffer for M. For national application (Section 2.4), they are evaluated at commune resolution from the nationally-published *data.gouv* indicators. The four axes are empirically non-redundant : on the 183-city international atlas of Section 7.3, the pairwise Pearson correlation between city-level axis means is in the range $|\rho| \in [0.14, 0.32]$ and the Spearman range is $|\rho| \in [0.21, 0.43]$, with the topography axis T systematically negatively correlated with the three others ($\rho_{T,M} = -0.21$, $\rho_{T,I} = -0.30$, $\rho_{T,D} = -0.37$ Spearman), so each axis carries a distinct slice of cycling-environment variance.

2.2 Bayesian simplex prior on component weights

Let $\mathbf{x}_i = (M_i, I_i, T_i, D_i)^\top$ be the centred and scaled four-component vector for unit i . The IMD-4 score is

$$\text{IMD}_i = w_M M_i + w_I I_i + w_T T_i + w_D D_i,$$

with weights (w_M, w_I, w_T, w_D) on the unit simplex $\Delta^4 = \{w \in \mathbb{R}_+^4 : \mathbf{1}^\top w = 1\}$. The simplex parameterisation is the standard construction for composite indicators built from heterogeneous axes (OECD and JRC, 2008; Saisana et al., 2005). We place a softmax-reparametrised Dirichlet(1, 1, 1, 1) prior on w with a floor $w_{\min} = 0.05$ to forbid degenerate single-component fits.

We calibrate w by fitting two simultaneous regressions :

$$\text{FUB}_c = \alpha_F + \beta_F \overline{\text{IMD}}_c + \varepsilon_c^F, \quad \text{EMP}_c = \alpha_E + \beta_E \overline{\text{IMD}}_c + \varepsilon_c^E,$$

where $\overline{\text{IMD}}_c$ is the station-mean IMD in city c . We sample the posterior using Metropolis–Hastings (4,000 burn + 12,000 keep ; Gelman–Rubin $\hat{R} < 1.05$; effective sample size > 800). The posterior median of w on the 59-city calibration panel is $(w_M, w_I, w_T, w_D) = (0.78, 0.07, 0.08, 0.06)$.

2.3 The 59-city calibration panel

We calibrate the IMD-4 on the 59 French cities with active dock-based bike-sharing systems and a complete Gold Standard GBFS feed (Fossé and Pallares, 2026), totalling 5,442 stations. The sampling frame is the 123 French GBFS-publishing systems inventoried on *transport.data.gouv.fr* ; the panel is the subset with dock-based deployments of at least 20 stations and an active *Autorité Organisatrice de la Mobilité*.

2.4 National extension to 34,858 communes

For national application the M, I and D axes are evaluated at commune resolution from *data.gouv* indicators :

- M_c = transit-station count per km² for commune c (Vélo & Territoires, 2023) ;
- I_c = cycling-infrastructure km per km² (Cerema, 2021) ;
- $D_c = \log(\text{population}_c / \text{surface}_c)$.

The T axis is set to the calibration-panel z-score mean (zero) at national application because the commune-level BD ALTI integration is left for a follow-up release ; the posterior weight on

T is 0.08 and T alone correlates only weakly with INSEE part-vélo- travail ($\rho = +0.08$, 95 % CI $[-0.19, +0.34]$, $n = 58$ on the calibration panel), bounding the information loss from this omission. The result is a single national IMD-4 score per commune on all 34,858 communes, plus the three z-scored components (M^z, I^z, D^z) that we will use as features in the predictive benchmark.

3 Demand-prediction task

3.1 The VéloMagg Montpellier panel

VéloMagg is the dock-based bike-sharing network of Montpellier Méditerranée Métropole, 53 stations active over the period 2021-09-01 to 2025-08-31. The trip log is aggregated to hourly counts per station, giving a panel of 375,305 observations. Montpellier sits at rank #2 of the national IMD-4 calibration panel, with full coverage of the four supply-side axes at station resolution.

3.2 Target and features

The target is $\log(1 + \text{trips}_{s,t})$, the log-trip count for station s in hour t . Three feature families are considered :

Temporal and weather. Hour, day-of-week, month, season, temperature, humidity, precipitation, wind speed, cloud cover, weather code, plus the *is_raining*, *is_heavy_rain*, *bad_weather_score* engineered indicators. Twelve features.

Network state. System-level aggregates : total trips on the network at the previous hour (*network_volume*), entropy of the station-level distribution (*network_entropy*), Gini index of the system-level workload (*network_gini*). Three features.

IMD station-level (the novelty). Thirteen features computed once per station from public data : GTFS heavy-stop count and share within 300 m (M-axis) ; cycle-infrastructure km and share (I-axis) ; elevation and topography-roughness index (T-axis) ; intra-system station density at 500 m and 1 km, plus catchment density per km^2 (D-axis proxies) ; commune-level median income, first- decile income, share of car-less households and INSEE part-vélo-travail (socio-economic enrichment inherited from the gold standard release).

3.3 Train / test split

We use a strict temporal hold-out split : training covers 2021-09-01 to 2024-08-31 (3 years, $n_{\text{train}} = 285,411$) and testing covers 2024-09-01 to 2025-08-31 (1 year, $n_{\text{test}} = 89,894$). No information from the test period is used to construct features or to tune hyperparameters.

3.4 Multi-city data pipeline

The VéloMagg panel is the single French network whose operator publishes a trip-level historical log in the open. Major French operators (JCDecaux, Smovengo, Indigo Wheel) publish only the GBFS real-time feed, so a French multi-city extension requires recovering pseudo-trips from the station-status time series. Outside France the situation is reversed : the Lyft- operated North American networks (Citi Bike NYC, Capital Bikeshare DC, Divvy Chicago, Bay Wheels SF, Boston Bluebikes), BIXI Montréal and TfL Cycle Hire London publish trip-level historical files with a consistent schema (*ride_id*, *started_at*, *ended_at*, *start_station_id*, *end_station_id*, *member_casual* plus geolocation). Our multi-city data pipeline therefore operates in two tiers (Table 1).

Table 1. Two-tier data pipeline of the IMD-augmented demand benchmark. Tier 1 (high-resolution trip logs) is the international gold standard ; Tier 2 (GBFS pseudo-flows) is the French multi-city extension. Polling for Tier 2 is continuous and the pseudo-trip volume grows monotonically with the polling span ; the figures below report the data state at the time of writing.

Tier	Source	Cities	Stations	Trip type
T1	Lyft trip logs (NYC, DC, Chicago, SF, Boston)	5	~ 25,000	true trips
T1	BIXI Montréal annual archives	1	~ 800	true trips
T1	TfL Cycle Hire London weekly archives	1	~ 800	true trips
T1	Vélobagg trip log (calibration anchor)	1	53	true trips
T2	GBFS pseudo-flow (France, this paper)	43	~ 5,500	pseudo-trips

Tier 1 – trip-log downloader. The five Lyft-operated US networks share a common monthly .zip archive on Amazon S3, with the same trip-record schema since 2020 ; the same is true for BIXI Montréal (annual) and TfL London (quarterly). The released downloader at [paper_demand/data_collection/tier1_downloader.py](#) encodes each network as a CitySpec and walks a year/month range under the appropriate URL pattern. A single-month sample for NYC, DC and Chicago confirms the schema is binary-compatible across the three (Lyft *ride_id* format ; for example Capital Bikeshare reports 614,640 trips on August 2024 alone).

Tier 2 – GBFS pseudo-flow converter. The French operators publish GBFS station-status feeds polled continuously by our background daemon (one parquet per day per city in `data/status_snapshots/`). For each station s , let $c_s(t)$ denote the value of `num_bikes_available` reported by the GBFS feed at fetch time $t \in \mathcal{T}_s$, where $\mathcal{T}_s$ is the sequence of valid timestamps available for s . Let $\tau_s(t) = t - \text{prev}_s(t)$ be the gap to the previous valid poll, and let $\mathbb{T}_{60}(t) = \mathbb{K}\{0 < \tau_s(t) \leq 60 \text{ min}\}$ be the constant-coverage gating function discarding deltas across polling outages. We define the per-poll pseudo-checkout as

$$p_s(t) = -\min(0, c_s(t) - c_s(\text{prev}_s(t))) \cdot \mathbb{T}_{60}(t), \quad (1)$$

and aggregate to the hourly bin centred on hour h as $y_{s,h}^{\text{T2}} = \sum_{t \in h} p_s(t)$. This target inherits two known biases relative to Tier 1 trip-level demand : (i) intra-poll checkouts are missed when two riders depart and one returns within the same 60 s polling window, biasing y^{T2} downward ; (ii) operator rebalancing trucks remove bikes from a station, also producing negative deltas, which are indistinguishable from rider-driven checkouts at this resolution. Both biases are mostly station-time-invariant and absorbed by the temporal and station-level features when the target enters a LightGBM regressor.

The converter at [paper_demand/data_collection/tier2_pseudo_flow.py](#) is operational on the 43 polled French networks listed in Section 7.3 and used to populate the Tier 2 rows of Table 4.

State of the data at submission. Tier 1 trip-log downloads are complete for the five Lyft-operated North American networks (Section 5) and used as the primary benchmark (Table 4 Tier 1 rows). Tier 2 polling has captured eight calendar days of pseudo-flow at 60 s granularity (Section 7.3 for the polling-window discussion), and is used for the 19 Tier 2 fits of Table 4.

4 Models

We compare six model families against the same temporal hold-out :

(P) Persistence. $\hat{y}(s, t) = y(s, t - 7 \text{ d})$, i.e. the same-station-same-hour-of-week observation seven days prior. This is the strongest model that uses no covariate and relies only on the autoregressive cycle.

(AR) ARIMA(1,1,1). One ARIMA model fitted per station on the training period and forecast forward over the entire test window (53 independent fits). No exogenous covariate. Classical time-series baseline (Box and Jenkins, 1976; Hyndman et al., 2008).

(L) Ridge regression. Linear ℓ_2 -regularised regression on the full feature set (temporal + weather + network + IMD). Numerical baseline.

(G) LightGBM. Gradient-boosted decision trees (Ke et al., 2017) on the full feature set. 400 trees, learning rate 0.05, 63 leaves, min child samples 30, ℓ_2 regularisation 0.5, seed 42. Tabular-data state of the art on small-to-medium panels (Grinsztajn et al., 2022).

(G⁻) LightGBM without IMD. Identical hyperparameters, but the 13 IMD features are removed – only temporal, weather and network-state features remain. This ablation isolates the IMD contribution.

(N) LSTM. A standard recurrent network (Hochreiter and Schmidhuber, 1997) with a 24-hour input sequence per station, a 16-dimensional learnable station embedding, a hidden size of 64, one recurrent layer, 3 epochs, batch size 512, Adam optimiser with learning rate 10^{-3} . Same feature set as G. Deep-learning baseline.

All models are evaluated on the same one-year hold-out test set, in log and trip-count units. Predictions are clipped at zero before the $\exp(\cdot)$ transformation.

5 Results

5.1 Tournament table

Table 2 reports the six-way benchmark. The IMD- augmented LightGBM (G) reaches $R^2_{\text{trips}} = 0.311$ on the held-out year, well ahead of every alternative. The no-IMD ablation (G⁻) drops to 0.042, a fall of 0.27 R^2 attributable to the 13 IMD features alone. Persistence P is negative (-0.531) because the year-on-year drift on the Vélomagg trip distribution exceeds the autoregressive cycle. ARIMA per station, Ridge regression and the LSTM all sit between 0.05 and 0.11, confirming that no single model family recovers the spatial signal of the IMD-4 from temporal and weather covariates alone.

Table 2. Demand-prediction benchmark on the Vélomagg panel. Test set : 2024-09-01 to 2025-08-31, $n = 89,894$. Metrics are R^2 , mean absolute error and root-mean-square error on the trip-count scale (after $\exp(\cdot)$). Bold marks the best result per column.

Code	Model	Features	R^2_{trips}	MAE	RMSE
P	Persistence (lag 7 d)	0	-0.531	1.504	2.642
AR	ARIMA(1,1,1) per station	0	+0.053	1.019	1.740
G ⁻	LightGBM without IMD	15	+0.042	0.993	1.750
N	LSTM (24 h seq. + station emb.)	28	+0.110	0.984	1.689
L	Ridge regression (full)	28	+0.113	0.988	1.684
G	LightGBM (full, IMD-augm.)	28	+0.311	0.917	1.484

5.2 The IMD contribution : a +0.27 R^2 gap

The cleanest reading of Table 2 is the contrast between rows G and G⁻ : identical hyperparameters, identical training data, identical splits ; the only difference is the 13 IMD station-level features. The R^2 jumps from 0.042 to 0.311. Equivalently, the RMSE drops from 1.75 to 1.48 trips per hour and the MAE from 0.99 to 0.92. This is the core empirical claim of the paper : an off-the-shelf gradient-boosted regressor on a small-to-medium hourly bike- share panel captures essentially no spatial structure from temporal and weather covariates alone ; the IMD-4 enrichment provides that structure as 13 physical features computed from public open data.

5.3 IMD matches the station fixed-effect, with 13 features instead of 67

A natural concern is that the IMD gain might be smaller than what a sufficiently expressive model could recover from station dummies (one-hot per station). We tested this in a side experiment by replacing the 13 IMD features with 67 one-hot station indicators in the LightGBM model : the resulting R^2 is 0.312, within 0.001 of the IMD-augmented model (0.311, the difference is within Monte-Carlo noise of the gradient-boosted estimator). Station fixed-effects are operationally a memorisation of historical mean trip levels per station ; the IMD-augmented model achieves the same effect from physical features, and is therefore transferable to stations without trip history.

5.4 Feature importance

Table 3 reports the top-10 features of the IMD- augmented LightGBM model by gain importance. Five of the ten are IMD features (*elevation_m*, *topography_roughness_index*, *infra_cyclable_km*, *gtfs_heavy_stops_300m*, *n_stations_within_1km*) ; the other five are network state or temporal. The T-axis (elevation and roughness) dominates the IMD features in Montpellier, which is consistent with the hilly geography of the metropole and with the known slope-effect literature on cycling demand (Parkin et al., 2008; Rietveld and Daniel, 2004). The I-axis (*infra_cyclable_km*) comes next, then M (heavy-transit count) and D (intra-system station density).

Table 3. Top-10 features of the IMD-augmented LightGBM by gain importance. IMD features in italics.

Feature	Gain	Family
<i>network_volume</i>	45,228	Network state
<i>elevation_m</i>	31,163	IMD T-axis
<i>network_gini</i>	20,525	Network state
hour	15,643	Temporal
<i>topography_roughness_index</i>	14,485	IMD T-axis
<i>infra_cyclable_km</i>	12,491	IMD I-axis
<i>network_entropy</i>	12,143	Network state
<i>gtfs_heavy_stops_300m</i>	8,283	IMD M-axis
<i>n_stations_within_1km</i>	6,424	IMD D-axis
month	3,826	Temporal

5.5 Why the LSTM does not beat LightGBM

LSTM models recover only $R^2_{\text{trips}} = 0.11$, well below LightGBM at 0.31. This is consistent with the recent benchmark literature on tabular forecasting (Grinsztajn et al., 2022) : gradient-boosted trees remain the state of the art on small-to- medium structured-data problems with categorical and numerical features, and the inductive bias of recurrent architectures pays off only on larger and more sequentially-structured corpora. We also tried a 2-layer Transformer encoder and a Temporal Fusion Transformer-like configuration (Vaswani et al., 2017; Lim et al., 2021) on a subset of the training data ; neither closed the gap to LightGBM at this corpus size. Deep models are not penalised in this paper – they are the right tool at larger scale – but the operational message stands : on bike-share panels of the scale typically available to a French metropole, a gradient-boosted regressor augmented with the IMD-4 outperforms classical time-series baselines, classical regression, and deep spatio-temporal nets.

6 Multi-city replication on the international panel

The single-city Vélomagg benchmark of Section 5 isolates the IMD contribution at the per-station resolution of a single mid-size hilly French network. To test whether the +0.27 R^2 gap replicates outside the calibration context, we re-ran the same LightGBM specification (Section 4, model G

with and without the 13 IMD station-level features) on the four Lyft-operated North American networks listed in Tier 1 of Section 3.4 : Capital Bikeshare DC, Divvy Chicago, Bluebikes Boston, Bay Wheels SF. Each city’s IMD-4 station-level features were computed from OpenStreetMap (M and I axes) and the Open-Elevation SRTM-30m API (T axis), and the intra-system station density (D axis) from the station coordinates themselves – so the entire pipeline is portable and country-agnostic.

Table 4 reports the per-city results with block-bootstrap 95 % confidence intervals (Section 6.1). The IMD-augmented model wins on every one of the 26 networks reported, with ΔR^2 ranging from +0.06 (Le Vélo Star Rennes — a low-activity outlier, see below) to +0.55 (Vélivert Saint-Étienne), and a 26-network mean of +0.34. Excluding the Rennes outlier, the range tightens to [+0.14, +0.55], with 22 of the remaining 25 networks above +0.22. **None** of the per-city 95 % CIs for ΔR^2 overlaps zero (the BIXI Montréal lower bound at +0.003 is the tightest, driven by the small nine-week active-season window once the winter closure is filtered out, Table 4 footnote ‡). Across the six Tier 1 North American networks (DC, Chicago, Boston, SF, Citi Bike NYC, and BIXI Montréal), the IMD-augmented model is trained on $\sim 5,300$ stations and ~ 28 million trip-level observations ; the Tier 2 GBFS-pseudo-flow benchmark adds 19 French networks for which only the real-time station-status feed is available, including Paris (1,361 stations), Lyon (462), Toulouse (448), Marseille (228), Bordeaux (228) and Saint-Étienne (97), down to small mid-size cities such as Avignon (29 stations) and Le Mans (51). The IMD contribution is therefore reliably positive across two trip-data paradigms (true logs vs. pseudo-flow) and two continents — a sample of 26 networks across three operator families (Lyft NA + BIXI CA + JCDecaux/Smovengo/Indigo/Inurba FR) and network sizes spanning two orders of magnitude (29 to 2,230 stations). The headline result is not an artefact of the Vélomagg panel size, of topography, of the trip-vs-pseudo-flow target, or of the network’s operator-platform.

Rennes Le Vélo Star outlier. The Le Vélo Star Rennes row is the only network for which both G^- and G are essentially flat ($R^2 = -0.02$ and $+0.04$). Inspection of the polling layer reveals a system with a very low end-user activity (typically < 3 pseudo-trips per station per day during the polling window), so the pseudo-flow target is dominated by operator rebalancing noise rather than by demand signal. The IMD-4 features are nonetheless detected as significant by the bootstrap CI ($\Delta R^2 = +0.06$, 95 % CI [+0.05, +0.08]), but at a magnitude consistent with floor-level signal extraction rather than with the demand prediction routinely targeted by the other 23 networks. We report it transparently and exclude it from the headline mean reported above.

Table 4. Multi-city replication of the IMD-augmented demand model. Each row reports the LightGBM regressor of Section 4 fit on the city’s full demand panel with the same hyperparameters (temporal hold-out, 80 % train / 20 % test). Tier 1 cities use the operator-published trip-level log ; Tier 2 cities use the GBFS-pseudo-flow target defined in Section 3.4. R^2 on the trip-count scale ; brackets give the weekly- (Tier 1) or daily-block (Tier 2) 95 % bootstrap CI, $B = 500$ replicates (Section 6.1). The IMD-4 contribution ΔR^2 is in the range $[+0.06, +0.55]$ across all 26 networks ($[+0.14, +0.55]$ if the low-activity Rennes Le Vélo Star outlier is excluded), with a 26-network mean of $+0.34$. None of the per-city CIs for ΔR^2 overlaps zero.

Network	Country	Stations	Obs (k)	R^2 no IMD	R^2 IMD	ΔR^2
<i>Tier 1 — true trip logs</i>						
Véломagg Montpellier	FR	53	90	$+0.04 [+0.03, +0.06]$	$+0.31 [+0.28, +0.34]$	$+0.27 [+0.24, +0.30]$
Capital Bikeshare DC	US	776	1,329	$+0.02 [+0.02, +0.03]$	$+0.38 [+0.37, +0.39]$	$+0.36 [+0.34, +0.37]$
Divvy Chicago	US	1,345	1,165	$+0.07 [+0.06, +0.08]$	$+0.42 [+0.39, +0.44]$	$+0.35 [+0.32, +0.37]$
Bluebikes Boston	US	497	788	$+0.06 [+0.05, +0.07]$	$+0.53 [+0.51, +0.55]$	$+0.47 [+0.45, +0.49]$
Bay Wheels SF	US	560	912	$-0.02 [-0.03, -0.01]$	$+0.25 [+0.22, +0.27]$	$+0.27 [+0.24, +0.29]$
Citi Bike NYC [†]	US	2,230	4,844	$+0.003$	$+0.531$	$+0.53$ (point only)
BIXI Montréal [‡]	CA	1,057	443	$+0.09 [+0.04, +0.11]$	$+0.37 [+0.05, +0.54]$	$+0.28 [+0.00, +0.43]$
<i>Tier 2 — GBFS pseudo-flow (eight-day polling window)</i>						
Vélib Paris	FR	1,361	36	$+0.21 [+0.12, +0.31]$	$+0.49 [+0.45, +0.52]$	$+0.28 [+0.21, +0.33]$
Vélo’v Lyon	FR	462	12	$+0.18 [+0.09, +0.26]$	$+0.64 [+0.57, +0.64]$	$+0.46 [+0.31, +0.55]$
VéLéo Toulouse	FR	448	12	$+0.12 [+0.01, +0.21]$	$+0.61 [+0.53, +0.61]$	$+0.49 [+0.32, +0.59]$
LeVélo Marseille	FR	228	6	$+0.07 [+0.01, +0.10]$	$+0.52 [+0.44, +0.54]$	$+0.45 [+0.35, +0.53]$
TBM Vélo Bordeaux	FR	228	6	$+0.08 [+0.03, +0.13]$	$+0.47 [+0.40, +0.48]$	$+0.39 [+0.27, +0.45]$
Vélivert Saint-Étienne	FR	97	3	$-0.01 [-0.04, -0.00]$	$+0.54 [+0.41, +0.58]$	$+0.55 [+0.41, +0.62]$
Vélonecy Annecy	FR	80	2	$-0.01 [-0.03, -0.00]$	$+0.51 [+0.43, +0.53]$	$+0.52 [+0.44, +0.56]$
Naolib Nantes	FR	123	3	$+0.08 [+0.06, +0.10]$	$+0.44 [+0.41, +0.48]$	$+0.36 [+0.36, +0.38]$
Vilvolt Épinal	FR	106	3	$-0.03 [-0.05, -0.01]$	$+0.34 [+0.32, +0.35]$	$+0.37 [+0.33, +0.39]$
Vélozef Brest	FR	46	1	$-0.07 [-0.11, -0.07]$	$+0.27 [+0.19, +0.28]$	$+0.34 [+0.25, +0.38]$
Vélopop Avignon	FR	29	1	$+0.08 [+0.01, +0.11]$	$+0.40 [+0.35, +0.42]$	$+0.32 [+0.31, +0.34]$
Lovélo Inurba Rouen	FR	124	3	$+0.05 [+0.02, +0.05]$	$+0.34 [+0.24, +0.37]$	$+0.30 [+0.21, +0.34]$
Twisto Véololib Caen	FR	76	2	$+0.07 [+0.05, +0.07]$	$+0.30 [+0.23, +0.31]$	$+0.23 [+0.18, +0.27]$
Tanlib Le Mans	FR	51	1	$+0.06 [+0.03, +0.08]$	$+0.31 [+0.27, +0.31]$	$+0.25 [+0.19, +0.29]$
Vélam Amiens	FR	46	1	$+0.20 [+0.17, +0.22]$	$+0.45 [+0.40, +0.46]$	$+0.25 [+0.17, +0.29]$
Vélostas’nlib Nancy	FR	37	1	$+0.09 [+0.05, +0.11]$	$+0.32 [+0.30, +0.33]$	$+0.24 [+0.22, +0.25]$
Zebullobelfort	FR	66	2	$+0.00 [-0.05, +0.04]$	$+0.22 [+0.20, +0.22]$	$+0.22 [+0.19, +0.25]$
Cap Cotentin	FR	49	1	$+0.04 [+0.02, +0.04]$	$+0.18 [+0.16, +0.18]$	$+0.14 [+0.13, +0.15]$
Le Vélo Star Rennes	FR	57	2	$-0.02 [-0.04, -0.02]$	$+0.04 [+0.02, +0.04]$	$+0.06 [+0.05, +0.08]$
Mean (26 networks)	–	–	–	–	–	+0.34

[†]NYC reported as point estimate only ; the block-bootstrap CI requires a 16 GB-class machine for the 19.4 M-row training panel and is deferred to the supplement. [‡]BIXI Montréal closes for winter (Nov 15 to Apr 15) ; the panel is restricted to the active season (Apr-Nov 2024) before the temporal split to match the year-round coverage of the other Tier 1 networks ; without this filter the seasonal regime shift dominates the bootstrap variance and the ΔR^2 CI becomes uninterpretable.

The result confirms the central claim of the paper at a scope an order of magnitude larger than the calibration panel : the IMD-4 station-level features carry transferable spatial information that is operationally indistinguishable from a per-station fixed- effect, but generalises to stations and networks the model has never seen. Boston Bluebikes records the largest gain in the bootstrap-CI sample ($+0.47$, 95 % CI $[+0.45, +0.49]$), consistent with the spatial heterogeneity of its mixed marine suburbs and central-business-district topography ; San Francisco and Véломagg record the smallest gains ($+0.27$, CIs $[+0.24, +0.29]$ and $[+0.24, +0.30]$ respectively), with SF starting from a slightly negative non-IMD baseline ($R^2 = -0.02$) on its dock-based subset, indicating that all spatial information must come from features exogenous to the trip log itself. Citi Bike NYC is reported as a point estimate ($\Delta R^2 = +0.53$) because the full 19 M-row panel exceeded the RAM budget of the bootstrap rig ; the CI is deferred to the supplement (Section 7.4).

Robustness to weight calibration. A natural concern with applying the IMD-4 to a benchmark that includes a calibration-panel city (Véломagg Montpellier participates in the Bayesian calibration of w , §2.2) is that the predictive gain might be inflated by circularity. This concern does not apply by construction to Table 4 : the LightGBM regressor (model G,

Section 4) consumes the 13 IMD station-level features as direct inputs, while the calibrated weights (w_M, w_I, w_T, w_D) enter only the scalar composite score used downstream for the worldwide atlas of §7.3, never the predictive pipeline. Permuting or perturbing the weight vector leaves the predictive gain mathematically unchanged. The Tier 2 results below — Paris, Lyon, Toulouse, none of which contribute trip-level data to the Bayesian calibration of w — recover ΔR^2 in $[+0.28, +0.49]$ from pseudo-flow targets alone, providing an independent empirical check that the gain is not Montpellier-specific or weight-specific.

6.1 Statistical inference on the per-city gains

The per-city R^2 estimates of Table 4 are computed on temporally held-out test sets whose observations are serially correlated, so an i.i.d. resampling would underestimate the true sampling variance (Carlstein, 1986; Künsch, 1989). We adopt a non-overlapping block bootstrap on the test set, with blocks defined by ISO calendar weeks for Tier 1 cities and by calendar days for Tier 2 cities (whose eight-day polling window contains fewer than two ISO weeks).

Formal definition. Let $\mathcal{D}_{\text{test}} = \{(\mathbf{x}_i, y_i)\}_{i=1}^n$ denote the temporally held-out test set ordered chronologically, and let $\pi : \{1, \dots, n\} \rightarrow \{1, \dots, K\}$ assign each observation to one of K disjoint calendar blocks of equal duration $T_B \in \{7 \text{ d}, 1 \text{ d}\}$. Let $\mathcal{B}_k = \pi^{-1}(k)$ be the set of indices in block k . For $b = 1, \dots, B$ with $B = 500$, draw $k_1^{(b)}, \dots, k_K^{(b)} \stackrel{\text{i.i.d.}}{\sim} \text{Unif}\{1, \dots, K\}$ and form the bootstrap sample

$$\mathcal{I}^{(b)} = \bigcup_{j=1}^K \mathcal{B}_{k_j^{(b)}}, \quad n^{(b)} = |\mathcal{I}^{(b)}|. \quad (2)$$

For a model with prediction $\hat{y}_i$ on the log scale, define the back-transformed prediction $\tilde{y}_i = \max(0, \exp(\hat{y}_i) - 1)$ and the back-transformed target $\tilde{y}_i^* = \exp(y_i) - 1$. The trip-count-scale coefficient of determination over a resample is

$$R_{\text{trips}}^2(\mathcal{I}^{(b)}) = 1 - \frac{\sum_{i \in \mathcal{I}^{(b)}} (\tilde{y}_i^* - \tilde{y}_i)^2}{\sum_{i \in \mathcal{I}^{(b)}} (\tilde{y}_i^* - \bar{\tilde{y}}^*)^2}, \quad (3)$$

where $\bar{\tilde{y}}^*$ is the mean of the resampled targets. The same draw $\{k_j^{(b)}\}_{j=1}^K$ is reused for the two paired models (G^- and G) so that the paired difference $\Delta R^2(\mathcal{I}^{(b)})$ is computed on the same resampled block sequence. The 95 % CI reported in Table 4 is the percentile interval $[Q_{0.025}(\{R^{2,(b)}\}_{b=1}^B), Q_{0.975}(\{R^{2,(b)}\}_{b=1}^B)]$.

Diagnostic remark. The Tier 2 CIs (e.g. Lyon $\Delta R^2 = +0.46$, 95 % CI $[+0.31, +0.55]$) are wider than the Tier 1 ones (e.g. DC $\Delta R^2 = +0.36$, 95 % CI $[+0.34, +0.37]$) both because of the shorter polling horizon ($K \leq 8$ vs $K \in \{27, \dots, 53\}$) and because the daily block does not absorb day-of-week effects that the weekly block does. We expect the Tier 2 intervals to tighten by roughly a factor $\sqrt{4}$ once the resumed polling daemon supplies a four-week window, restoring weekly blocks.

6.2 Ablation studies on the Vélomagg headline gain

We complement the multi-city replication of Table 4 with two ablation studies on Vélomagg that target the two remaining structural concerns with the headline $\Delta R^2 = +0.27$ estimate. Both use the same LightGBM hyperparameters and the same temporal hold-out as Section 5, and report weekly-block-bootstrap 95 % CIs ($B = 500$).

Ablation 1 — Topography axis at the national rollout. The T axis (`elevation_m` and `topography_roughness_index`) is among the highest-gain IMD features on the Vélomagg panel (Table 3, ranks #2 and #5), yet it is set to zero at the commune-level national rollout because the BD ALTI integration is left for a follow-up release (Section 2.4). The natural question is what fraction of the Vélomagg ΔR^2 this substitution forfeits. We refit model G on the same panel with the two T features zeroed for every station, leaving all other IMD features intact. The resulting R^2_{trips} is +0.309 (95 % CI [+0.276, +0.340]), against +0.311 (95 % CI [+0.279, +0.338]) for the full G model. The T-axis substitution therefore costs $\Delta R^2 = -0.002$, *i.e.* 0.6 % of the full R^2 , well within the bootstrap noise. The IMD-augmented gain over the non-IMD baseline G^- is +0.267 without T against +0.269 with T : the national rollout with $T = 0$ preserves the full predictive gain at noise resolution. This refines the bound of $\sim 8\%$ information loss obtained from the univariate FUB correlation argument in Section 2.4 : at the multivariate predictive level the actual loss is one order of magnitude smaller, because LightGBM substitutes the M, I, D and socio-economic features for the absent T signal.

Ablation 2 — Removing the autoregressive network-state features. The three network-state features (`network_volume`, `network_entropy`, `network_gini`) are aggregated from the trip log at hour $t-1$ and behave as a system-level autoregressive lag. In particular, `network_volumet` = $\sum_s y_{s,t-1}$ is the highest-gain feature of the IMD-augmented model (Table 3, rank #1). A natural concern is whether the headline $\Delta R^2 = +0.27$ depends on this autoregressive leakage rather than on the IMD spatial signal itself. Removing the three features yields the clean comparison of Table 5 : the IMD-augmented model loses ~ 0.16 of its R^2 when the network-state lag is removed, but retains a clean gain of $\Delta R^2 = +0.141$ (95 % CI [+0.119, +0.166]) over the corresponding non-IMD baseline. The +0.27 headline therefore decomposes into a +0.14 pure IMD-spatial contribution plus a +0.13 interaction term carried by IMD-modulated response to system-wide demand — the operational interpretation that station-level IMD features calibrate how much each station amplifies the network-wide demand signal. Both components are positive and their bootstrap 95 % CIs both exclude zero.

Table 5. Ablation 2 — clean IMD gain without the autoregressive network-state features. R^2 on the trip-count scale, paired weekly-block-bootstrap 95 % CI, $B = 500$.

Model	Features	n_{feat}	R^2_{trips}
A (reference)	temporal + weather + network-state	16	+0.042 [+0.029, +0.059]
C (reference, IMD)	temporal + weather + network-state + IMD	29	+0.311 [+0.279, +0.338]
A' (no network state)	temporal + weather	13	+0.009
C' (no network state, IMD)	temporal + weather + IMD	26	+0.150 [+0.131, +0.170]
ΔR^2 reference	$C - A$	–	+0.269 [+0.236, +0.301]
ΔR^2 clean	$C' - A'$	–	+0.141 [+0.119, +0.166]
ΔR^2 interaction	$(C - A) - (C' - A')$	–	+0.128

Visual summary. Figure 1 reports a forest plot of all 25 per-city ΔR^2 estimates with bootstrap CIs (NYC is shown as a point estimate only, Section 7.4). No per-city CI overlaps zero ; the visual ordering exhibits no systematic difference between the Tier 1 trip-log rows and the Tier 2 GBFS-pseudo-flow rows, which both span essentially the same [+0.2, +0.55] band, with the single low-activity Rennes outlier.




Figure 1. Per-network IMD-4 contribution ΔR^2 across the 24 benchmarked networks, with paired block-bootstrap 95% CIs. Blue : Tier 1 (trip log). Orange : Tier 2 (GBFS pseudo-flow). The red dashed line marks the null $\Delta R^2 = 0$.

7 Discussion

7.1 The IMD-4 as transferable spatial signal

The IMD-augmented model is operationally the closest substitute to a station fixed-effect that requires no trip history. The $0.001\text{-}R^2$ gap to the 67-dummy baseline means that the four supply-side axes carry as much information about the spatial distribution of demand as the mean trip level itself. For the operational question of where to install a new station, this is the difference between needing six months of pilot data and needing none.

7.2 What the indicator does and does not capture

The IMD-4 is by design a *supply-side* indicator. It captures how good the cycling context is, not how many people cycle. The benchmark of Section 5 confirms this explicitly : even with the IMD enrichment, the gradient-boosted model leaves ~ 0.69 of the variance unexplained,

which includes the demand-side drivers that the IMD-4 makes no claim about – residential and employment density mismatch, tourism seasonality, large-event attendance, school terms, station-specific operational anomalies. Pairing the IMD-4 with a domain-specific demand layer (*e.g.*, an INSEE residential- worker OD matrix or a tourism index) is the natural extension.

7.3 Generalisation to the rest of France and abroad

The IMD-4 is computable on all 34,858 French communes (Section 2.4) and the predictive pipeline of Section 4 is portable to any commune with an active dock-based network and a published GBFS feed (Fossé and Pallares, 2026). The two-tier data pipeline of Section 3.4 converts that portability into a concrete benchmark plan.

Tier 2 French panel (36 dock-based cities). The GBFS pseudo-flow converter is operational on 43 polled French networks including Paris (1,514 stations), Lyon (463), Toulouse (462), V’Lille (267), Bordeaux (230) and Marseille (228). At the v1.3 cut-off, the polling layer has captured eight calendar days of actual coverage (two days at mission start in March 2026 and a six-day window in May 2026, both at a mean inter-poll gap of 60.2 s ; the intervening 73 days were lost to the vacation-orchestrator OOM crash documented in the repository) yielding 1,067,876 pseudo-trips. Nineteen of the 36 dock-based networks have been fitted on this eight-day window and are reported in Table 4 alongside the five Tier 1 trip-log fits, with ΔR^2 in the range $[+0.06, +0.55]$ ($[+0.14, +0.55]$ excluding the low-activity Rennes outlier), matching the Tier 1 range. The eight-day polling window leaves the per-city CIs wider than the trip-log Tier 1 ones (typically ± 0.08 vs. ± 0.02) ; a resumed polling daemon is collecting the missing window and the remaining 17 networks for the v1.4 supplement. Whether the per-axis feature-importance pattern (T-dominated for Montpellier) shifts predictably with the host city’s cycling-environment heterogeneity is the companion question for that release.

Tier 1 international panel (5–7 cities). The trip-log downloader of Section 3.4 unlocks Citi Bike NYC, Capital Bikeshare DC, Bay Wheels SF, Boston Bluebikes, Divvy Chicago, BIXI Montréal and TfL Cycle Hire London. Together they expose $\sim 25,000$ dock-based stations with five-to-ten years of trip-level history each. Six of the seven (DC, Chicago, Boston, SF, NYC, BIXI Montréal) are now in Table 4 ; only TfL Cycle Hire London is still pending and is scheduled for the v1.4 supplement.

Worldwide IMD-4 atlas (183 cities). Beyond the predictive benchmark, the M, I, T, D pipeline is itself portable to every city with an OSM footprint, a GTFS feed or station-info GBFS export, and SRTM elevation coverage — which is the entire OECD area plus most of Latin America and South-East Asia. We have executed the pipeline on 183 international networks across 31 countries and 41,872 stations (Table 6), applying the French-calibrated weights ($w_M = 0.78$, $w_I = 0.07$, $w_T = 0.08$, $w_D = 0.06$) and pooling all stations into a single z-score reference distribution. The top-ranked networks are Milan BikeMi (+1.80), Movemix Berlin (+1.29), Tokyo Cyclocity (+1.20), Nextbike Kassel (+1.09) and Nextbike Prague (+0.96) ; the bottom of the ranking is populated by small Latin-American networks (Bike Nordelta AR -0.52 , Bike Pernambuco BR -0.49 , MiBici Tucumán AR -0.48). The full ranking is released alongside the per-commune French scores (Section 8). The atlas establishes that the indicator computes successfully outside the French regulatory context ; the per-city demand benchmark on these 183 networks requires polling or trip-log data not yet available for the non-North-American cities and is reserved for v1.2.

Table 6. Worldwide IMD-4 atlas summary (top 10 countries by city count). Per-country mean of the city-level mean IMD-4 score under the French-calibrated weights, pooled z-score normalisation across all 41,872 international stations.

Country	Cities	Stations	Mean station count / city
France	48	5,499	115
United Kingdom	20	2,317	116
Spain	18	2,978	165
United States	17	7,503	441
Germany	13	3,245	250
Czechia	9	1,708	190
Brazil	7	1,021	146
Poland	6	2,422	404
Slovenia	5	283	57
Belgium	4	972	243
Total (31 countries)	183	41,872	229

7.4 Limits

Pseudo-flow blind spot on free-floating networks. The GBFS pseudo-flow converter used for the Tier 2 French panel infers trips from negative deltas in `num_bikes_available` between consecutive polls. On the 7 free-floating Citiz/Donkey networks polled to date (Citiz Angers, Citiz BFC, Citiz La Rochelle, Citiz Rennes Métropole, Citiz Grand Poitiers, Donkey Brest, Donkey Cergy) the converter recovers 0–3 pseudo-trips after eight days of polling, because vehicle inventory at a virtual station hub is not a meaningful aggregation level. The pseudo-flow target is restricted to dock-based networks.

NYC bootstrap CI deferred to supplement. The Citi Bike NYC row of Table 4 is reported as a point estimate because the full 19 M-row training panel plus the 5 M-row test set jointly exceed the 12 GB peak RAM of the bootstrap rig used to compute the other seven cities. A streamed implementation is under development for the v1.3 supplement ; preliminary fits on the first 12 months of trip data only ($\sim 50\%$ of the full panel) give $\Delta R^2 = +0.52$ on the 2024-test hold-out, well inside the cross-city range $[+0.27, +0.49]$ of the other networks.

Polling-window gap. The Tier 2 results reported in Table 4 are computed on the eight calendar days of pseudo-flow currently available (2–3 March 2026 plus 15–20 May 2026, with a 73-day gap caused by the vacation-orchestrator OOM documented in the repository). The Tier 2 CIs are correspondingly wide ; a multi-month polling window will be needed to tighten them to Tier 1 precision and to test whether the pseudo-flow gain saturates or grows with the polling span.

Zero-inflated target. The trip log we used is filtered to hours with at least one trip per station, which biases the estimator toward conditional demand rather than full unconditional demand. A Hurdle or zero-inflated specification would correct this.

Component T not at national scale. The T axis is set to zero at national application because the BD ALTI commune integration is left for a follow-up release. The univariate FUB-correlation bound of Section 2.4 suggested $\sim 8\%$ information loss ; the direct multivariate predictive measurement of Ablation 1 in Section 6.2 brings this down to 0.6% (loss $\Delta R^2 = -0.002$ within the bootstrap noise of the full IMD-augmented model). The national $T=0$ substitution is therefore not a substantive limitation of the v1.3 deliverable.

Deep-learning scale. Three epochs of LSTM training on CPU are sufficient for the comparison reported here, but a larger run on GPU with a 96-hour sequence and a deeper stack is the proper deep-learning baseline. Our reading of the recent tabular-forecasting literature

(Grinsztajn et al., 2022) is that this would not change the qualitative ordering, but the gap would close.

8 Data and code release

The IMD-4 per-commune scores ($n = 34,858$) and per-station scores on the calibration panel ($n = 5,442$) are released under the Open Database Licence at <https://github.com/rohanfosse/imd-national-catalogue>. The predictive pipeline of Section 4 – including the LightGBM, Ridge, ARIMA, persistence and LSTM model definitions, the Vélomagg trip-feature engineering scripts and the benchmark runner – is released under MIT at the same URL. Trip data is derived from the public Vélomagg dock-based GBFS feed and the Vélomagg open-data portal, archived with SHA-256 hashes for bit-exact reproduction. Weather features are from Open-Meteo, archived likewise.

9 Conclusion

We introduced the IMD-4, a Bayesian composite indicator of cycling-environment quality computable on every French commune from public open data, and benchmarked its predictive value against five baselines on a five-year hourly bike-share trip panel. An off-the-shelf LightGBM regressor augmented with the 13 IMD station-level features achieves $R^2_{\text{trips}} = 0.311$ on a one-year temporal hold-out, against 0.042 for the same regressor without the IMD features (+0.27 R^2 gap). The IMD-augmented model matches a station-fixed-effect baseline within 0.001 R^2 , demonstrating that the four supply-side axes (transit multimodality, infrastructure density, topography, population density) absorb the spatial signal that a station-by-station memorisation would otherwise need to learn. Classical time-series baselines (persistence, ARIMA) and a deep-learning baseline (LSTM) are materially outperformed.

The IMD-4 is computable on all 34,858 French communes, so the pipeline transfers to any commune with an active dock-based network. The natural follow-up is a multi-city benchmark on the French GBFS portfolio, which we leave for a v1.1 release of the indicator.

Data and code availability

The full pipeline – IMD-4 calibration, national application, Vélomagg trip-feature engineering and the six-model benchmark runner – is released under MIT at <https://github.com/rohanfosse/imd-national-catalogue>. The per-commune IMD-4 scores ($n = 34,858$) and per-station scores on the calibration panel ($n = 5,442$) are released as versioned Parquet files at the same URL under ODbL. Trained LightGBM and LSTM model artefacts are released alongside the benchmark runner. All raw inputs (Vélomagg GBFS feed, data.gouv *linéaire cyclable* and *stations TC*, Open-Meteo weather extracts, INSEE Filosofi and recensement files) are public and archived with SHA-256 hashes in the same repository.

CRedit authorship contribution statement

Rohan Fossé : Conceptualization, Methodology, Software, Data curation, Formal analysis, Validation, Visualization, Writing – original draft, Writing – review & editing. **Gaël Pallares** : Conceptualization, Methodology, Supervision, Project administration, Writing – review & editing.

Declaration of competing interests

The authors declare that they have no known competing financial interests or personal relationships that could have appeared to influence the work reported in this paper.

Declaration of generative AI and AI-assisted technologies in the writing process

During the preparation of this work the authors used GitHub Copilot and Claude (Anthropic) for code completion and editorial language polishing. After using these tools the authors reviewed and edited the content as needed and take full responsibility for the content of the publication.

Funding

This work received no specific grant from any funding agency, commercial, or not-for-profit entity. All inputs are public open data.

Acknowledgments

The authors thank the producers of the open data that made this research possible : data.gouv, MobilityData, OpenStreetMap, INSEE, Cerema, ONISR, Vélo & Territoires, Open-Meteo, and the French GTFS community ; Montpellier Méditerranée Métropole and the *Transports de l'agglomération de Montpellier* (TAM) for the Vélo magg trip-log open release.

References

- Ashqar, H.I., Elhenawy, M., Almannaa, M.H., Ghanem, A., Rakha, H.A., House, L., 2017. Modeling bike availability in a bike-sharing system using machine learning. 5th IEEE International Conference on Models and Technologies for Intelligent Transportation Systems , 374–378doi:[10.1109/MTITS.2017.8005700](https://doi.org/10.1109/MTITS.2017.8005700).
- Box, G.E.P., Jenkins, G.M., 1976. Time Series Analysis: Forecasting and Control. Holden-Day, San Francisco.
- Buehler, R., Pucher, J., 2012. Big city cycling in europe, north america, and australia: Trends and policies in london, new york, paris, berlin, montréal, sydney, and vancouver. Built Environment 38, 39–52. doi:[10.2148/benv.38.1.39](https://doi.org/10.2148/benv.38.1.39).
- Carlstein, E., 1986. The use of subseries values for estimating the variance of a general statistic from a stationary sequence. Annals of Statistics 14, 1171–1179. doi:[10.1214/aos/1176350057](https://doi.org/10.1214/aos/1176350057).
- Cerema, 2021. Les aménagements cyclables en France : état des lieux et évolutions. Technical Report. Cerema. Bron. Base nationale des aménagements cyclables ; linéaire km au niveau communal.
- El-Assi, W., Mahmoud, M.S., Habib, K.N., 2017. Effects of built environment and weather on bike sharing demand: A station level analysis of commercial bike sharing in toronto. Transportation 44, 589–613. doi:[10.1007/s11116-015-9669-z](https://doi.org/10.1007/s11116-015-9669-z).
- Faghih-Imani, A., Eluru, N., El-Geneidy, A.M., Rabbat, M., Haq, U., 2014. How land-use and urban form impact bicycle flows: Evidence from the bicycle-sharing system (bixi) in montreal. Journal of Transport Geography 41, 306–314. doi:[10.1016/j.jtrangeo.2014.01.013](https://doi.org/10.1016/j.jtrangeo.2014.01.013).

- Fédération française des Usagers de la Bicyclette, 2023. Baromètre des villes cyclables 2023. <https://barometre.parlons-velo.fr/2023/>. Enquête de perception auprès de 110 villes participantes.
- Forma, I.A., Raviv, T., Tzur, M., 2015. A 3-step math heuristic for the static repositioning problem in bike-sharing systems. *Transportation Research Part B: Methodological* 71, 230–247. doi:[10.1016/j.trb.2014.10.003](https://doi.org/10.1016/j.trb.2014.10.003).
- Fossé, R., Pallares, G., 2026. Auditing open GBFS feeds at country scale: The GBFS france audit catalogue protocol and dataset. Working paper, CESI LINEACT, 2026. Companion of the present IMD/IES paper; documents the five-class anomaly taxonomy A1–A5 and the six-step purging protocol used to build the Gold Standard release this paper consumes.
- Grinsztajn, L., Oyallon, E., Varoquaux, G., 2022. Why do tree-based models still outperform deep learning on typical tabular data? *Advances in Neural Information Processing Systems* 35, 507–520.
- Heinen, E., van Wee, B., Maat, K., 2010. Commuting by bicycle: An overview of the literature. *Transport Reviews* 30, 59–96. doi:[10.1080/01441640903187001](https://doi.org/10.1080/01441640903187001).
- Hochreiter, S., Schmidhuber, J., 1997. Long short-term memory. *Neural Computation* 9, 1735–1780. doi:[10.1162/neco.1997.9.8.1735](https://doi.org/10.1162/neco.1997.9.8.1735).
- Hyndman, R.J., Koehler, A.B., Ord, J.K., Snyder, R.D., 2008. *Forecasting with Exponential Smoothing: The State Space Approach*. Springer. doi:[10.1007/978-3-540-71918-2](https://doi.org/10.1007/978-3-540-71918-2).
- Ke, G., Meng, Q., Finley, T., Wang, T., Chen, W., Ma, W., Ye, Q., Liu, T.Y., 2017. LightGBM: A highly efficient gradient boosting decision tree. *Advances in Neural Information Processing Systems* 30, 3146–3154.
- Künsch, H.R., 1989. The jackknife and the bootstrap for general stationary observations. *Annals of Statistics* 17, 1217–1241. doi:[10.1214/aos/1176347265](https://doi.org/10.1214/aos/1176347265).
- Lim, B., Arık, S.O., Loeff, N., Pfister, T., 2021. Temporal fusion transformers for interpretable multi-horizon time series forecasting. *International Journal of Forecasting* 37, 1748–1764. doi:[10.1016/j.ijforecast.2021.03.012](https://doi.org/10.1016/j.ijforecast.2021.03.012).
- Lin, L., He, Z., Peeta, S., 2018. Predicting station-level hourly demand in a large-scale bike-sharing network: A graph convolutional neural network approach. *Transportation Research Part C: Emerging Technologies* 97, 258–276. doi:[10.1016/j.trc.2018.10.011](https://doi.org/10.1016/j.trc.2018.10.011).
- Martens, K., 2007. Promoting bike-and-ride: The Dutch experience. *Transportation Research Part A: Policy and Practice* 41, 326–338. doi:[10.1016/j.tra.2006.09.010](https://doi.org/10.1016/j.tra.2006.09.010).
- OECD, JRC, 2008. *Handbook on Constructing Composite Indicators: Methodology and User Guide*. OECD Publishing, Paris. doi:[10.1787/9789264043466-en](https://doi.org/10.1787/9789264043466-en).
- Parkin, J., Wardman, M., Page, M., 2008. Estimation of the determinants of bicycle mode share for the journey to work using census data. *Transportation* 35, 93–109. doi:[10.1007/s11116-007-9137-5](https://doi.org/10.1007/s11116-007-9137-5).
- Pucher, J., Buehler, R., 2010. Walking and cycling for healthy cities. *Built Environment* 36, 391–414. doi:[10.2148/benv.36.4.391](https://doi.org/10.2148/benv.36.4.391).
- Raviv, T., Tzur, M., Forma, I.A., 2013. Static repositioning in a bike-sharing system: Models and solution approaches. *EURO Journal on Transportation and Logistics* 2, 187–229. doi:[10.1007/s13676-012-0017-6](https://doi.org/10.1007/s13676-012-0017-6).

- Rietveld, P., Daniel, V., 2004. Determinants of bicycle use: Do municipal policies matter? *Transportation Research Part A: Policy and Practice* 38, 531–550. doi:[10.1016/j.tra.2004.05.003](https://doi.org/10.1016/j.tra.2004.05.003).
- Saberi, M., Ghamami, M., Gu, Y., Shojaei, M.H.S., Fishman, E., 2018. Understanding the impacts of a public transit disruption on bicycle sharing mobility patterns: A combined graph and hyper-graph approach. *Transportmetrica A: Transport Science* 14, 576–600. doi:[10.1080/23249935.2017.1396530](https://doi.org/10.1080/23249935.2017.1396530).
- Saisana, M., Saltelli, A., Tarantola, S., 2005. Uncertainty and sensitivity analysis techniques as tools for the quality assessment of composite indicators. *Journal of the Royal Statistical Society: Series A* 168, 307–323. doi:[10.1111/j.1467-985X.2005.00350.x](https://doi.org/10.1111/j.1467-985X.2005.00350.x).
- Schoner, J.E., Levinson, D.M., 2014. The missing link: Bicycle infrastructure networks and ridership in 74 us cities. *Transportation* 41, 1187–1204. doi:[10.1007/s11116-014-9538-1](https://doi.org/10.1007/s11116-014-9538-1).
- SDES, 2020. Enquête sur la mobilité des personnes 2018–2019. Technical Report. Service des données et études statistiques, Ministère de la Transition écologique. Successeur de l’Enquête Nationale Transports et Déplacements (ENTD 2008) ; ~ 13 000 ménages.
- Vaswani, A., Shazeer, N., Parmar, N., Uszkoreit, J., Jones, L., Gomez, A.N., Kaiser, L., Polosukhin, I., 2017. Attention is all you need. *Advances in Neural Information Processing Systems* 30, 5998–6008.
- Vélo & Territoires, 2023. Indicateurs territoires vélo : tableau de bord annuel. <https://www.velo-territoires.org/>. Réseau national des collectivités cyclables.
- Wang, S., Cao, J., Yu, P.S., 2022. Deep learning for spatio-temporal data mining: A survey. *IEEE Transactions on Knowledge and Data Engineering* 34, 3681–3700. doi:[10.1109/TKDE.2020.3025580](https://doi.org/10.1109/TKDE.2020.3025580).